

Last time

- first steps with MPI

Today

- Computational scaling

→ how is my code performing in parallel?

60s → 90s → 2010s thing
Strong, weak, machine scaling.

- Data decomposition choices. →

Vectors.

Strong scaling

$$S_p = \frac{T_1}{T_p}$$

← true on one process
← on p processes.
speedup

Problem size is fixed.

Hope $S_p > 1$ when $p > 1$.

1 person can build a wall in

1 week. How long with 4

people take? $T_4 = \frac{T_1}{4}$ → ideal linear scaling

Gene Amdahl parallel part.

$$T_p = f T_1 + (1-f) \frac{T_1}{p}$$

serial fraction

parallel fraction

$$f T_1$$

$$(1-f) T_1$$

Amdahl's law:

$$\lim_{p \rightarrow \infty} T_p = f T_1$$

$$\Leftrightarrow \lim_{p \rightarrow \infty} S_p = \frac{1}{f}$$

So, if 1% of the code is serial, best speedup is?

$$S_\infty = \frac{1}{0.01} = 100.$$


Consequence

→ Need to strip out as much of the serial code as possible.

Presently speedup data:

Problem size N : ← fix.

Run varying P .



Note: dividing at by T_1 ,
so can hide inefficient implment.

Can also plot efficiency.

$\in [0, 1]$


$$\eta_p = \frac{T_1}{p T_p}$$

$p T_p$

total time we used on p processes.

exercise

What is η_p as a function of the serial fraction, f ?



processes

Often on log scales

Weak scaling


small, grows slowly.

$$T_p = T_1 + o(p)T_1$$

$$S_p = \frac{pT_1}{T_p}$$

~ Gustafson's law.
time to solve local problem.

local (per process) problem
size fixed.



Real applications scale problems
up with weak scaling \approx
then at fixed size do strong scaling.

$$\eta_P = \frac{T_1}{T_1 + o(p)T_1} = \frac{T_1}{1 + o(p)}$$

overhead fixed

$$\eta_P^{\text{fix}} = \frac{T_1}{1 + \alpha} ; \quad \zeta_P^{\text{fix}} = p \frac{1}{1 + \alpha}$$

$$\eta_P^{\log} = \frac{T_1}{1 + \log p} ; \quad \zeta_P^{\log} = \frac{p}{1 + \log p}$$

Typical setup for PDEs.

overhead: $\alpha + O(\log p)$

message latency

reductions

(combining data)

e.g. dot product $a \cdot b \rightarrow \text{has } \rightarrow \text{scale}$

"problem size fixed".

Dense matrix-matrix.

$N \times N$
 \mathbb{R}

$O(N^3)$ algorithm.

work scale.

If I double N .

~~to~~ double P
 \rightarrow the time to solve
fixed?

$$\begin{array}{l} N \rightarrow 2N \\ N^3 \quad \quad \quad 8N^3 \end{array} \quad , \quad \begin{array}{l} P \rightarrow 2P \\ \frac{N^3}{P} \rightarrow \frac{8N^3}{2P} \\ \quad \quad \quad = 4N^3 \end{array}$$

So to get constant T_p
need to add $8x$ processes

$$N^3 \rightarrow 8N^3$$

$$P \rightarrow 8P$$

$$\frac{N^3}{P} \rightarrow \frac{N^3}{P} \quad \checkmark$$

But local matrix size.

$$N \times N \rightarrow \frac{N}{\sqrt{8}} \times \frac{N}{\sqrt{8}} ?$$

so local problem
is smaller

\Rightarrow strong scaling.

Machine / algorithmic scale's

Range of sizes, Time to sol
plot T vs $\frac{N}{T} \rightarrow$ cost flat.

\rightarrow Add link to paper
w notes.

Vectors $\hat{=}$ parallel.

$$x \in \mathbb{R}^N$$

pointwise

$$y \leftarrow \alpha x + y$$

$$y \in \mathbb{R}^N$$

$$\alpha \in \mathbb{R}$$

$$y \leftarrow \frac{1}{y}$$

\vdots

collective.

computing norms/
dot products.

$$\beta \leftarrow a \cdot b$$

$$\beta \in \mathbb{R}$$


$$a, b \in \mathbb{R}^N$$

pointwise updates.

linearly we have
3 processes.

Assign to evenly distribute
total length N
across the p processes.

Scalability limit for
pointwise comes from
uneven distⁿ of
local work.



⇒ "load imbalance"

3, 3, 4

Dist products & grids
next time.